

M.TECH. PRACTICAL TRAINING REPORT
ON

A Literature Survey on Advanced Methodologies for
Vulnerability Detection in RESTful API Security

By

Mohd Gulam Mohd Hasim Ansari

Roll No :(242CS023)

Under the Guidance of

Professor P. Santhi Thilagam

Department of Computer Science & Engineering
NITK, Surathkal



DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING
NATIONAL INSTITUTE OF TECHNOLOGY KARNATAKA,
SURATHKAL, MANGLORE - 575025

DECLARATION

I hereby declare that the M. Tech. 3rd Semester **Practical Training** Report entitled **A Literature Survey on Advanced Methodologies for Vulnerability Detection in RESTful API Security**, which is being submitted to the National Institute of Technology Karnataka Surathkal, in partial fulfilment of the requirements for the award of the Degree of **Master of Technology in Computer Science and Engineering** in the Department of **Computer Science and Engineering**, is a bonafide report of the work carried out by me. The material contained in this Report has not been submitted to any University or Institution for the award of any degree.

Mohd Gulam Mohd Hasim Ansari

242CS023

Department of Computer Science and Engineering
NITK, Surathkal

Place: NITK, Surathkal.

Date:

CERTIFICATE

This is to certify that the M. Tech. 3rd Semester **Practical Training** Report entitled **A Literature Survey on Advanced Methodologies for Vulnerability Detection in RESTful API Security** submitted by **Mohd Gulam Ansari**, (Roll Number: 242CS023) as the record of the work carried out by him, is accepted as the M. Tech. 3rd sem report submission in partial fulfilment of the requirements for the award of the degree of **Master of Technology in Computer Science and Engineering** in the Department of **Computer Science and Engineering**.

Guide

Professor P. Santhi Thilagam

Department of Computer Science and Engineering

NITK, Surathkal

Chairman - DPGC

Dr. Manu Basavaraju

Department of Computer Science and Engineering

NITK, Surathkal

Contents

1	INTRODUCTION	1
2	REST API Architecture	4
2.1	The Client-Server Model	4
2.2	Communication Flow and Core Principles	5
2.3	The Role of the API Gateway	6
3	Authentication and Authorization	8
3.1	Session-Based Authentication: The Traditional, Stateful Approach	8
3.2	JWT-Based Authentication: The Modern, Stateless Approach	9
3.3	Summary of Authentication Methods	9
4	Case Study: Patient Management API Implementation	11
4.1	Technology Stack	11
4.2	API Endpoint Specification	12
4.3	Key Features and Functionality	14
4.4	Deployment and Live Demonstration	15
5	Common API Security Threats	16
5.1	Denial of Service Attacks	16
5.2	The OWASP API Security Top 10	16
6	LITERATURE SURVEY	18
6.1	Survey Table	18
7	RESEARCH GAPS	22
7.1	Identified Research Gaps	22

7.2 Strategies for Addressing the Gaps	22
8 Conclusion	24
9 Future Work	25
References	26

Chapter 1

INTRODUCTION

In our interconnected digital world, Application Programming Interfaces (APIs) are the invisible threads that weave our software together. At its core, an **API** is a set of rules and protocols that allows one software application to communicate with another, acting as a messenger that delivers requests and returns responses. This allows developers to access specific functionalities or data from other services without needing to understand their complex internal code, much like a diner uses a menu to order from a restaurant's kitchen without needing to know the recipes.

At the heart of this digital ecosystem is the **RESTful API**. The term REST, an acronym for **REpresentational State Transfer**, was first defined by computer scientist Roy Fielding in his 2000 doctoral dissertation (7). It is not a strict protocol but an architectural style that provides guiding principles for building scalable and efficient web services.[1] A RESTful API uses the standard HTTP protocol to perform actions on **resources**—any piece of information the API can provide, such as a user profile or a product listing. These actions, known as CRUD (Create, Read, Update, Delete) operations, are mapped directly to common HTTP methods: POST is used to create a resource, GET to read it, PUT to update it, and DELETE to remove it.

This architecture is built on several key principles. The first is a **client-server architecture**, which decouples the user interface (the client) from the data storage and logic (the server), allowing them to be developed and scaled independently.[1] The second, and perhaps most critical, principle is **statelessness**. This means that the server does not store any information about the client's session between requests; every request sent from the client must contain all the information needed for the server to process it.[1] This

lack of server-side session storage makes RESTful APIs highly reliable and scalable, as any server can handle any request at any time.

APIs can be broadly categorized by their intended audience and access policies. **Public APIs**, also known as Open APIs, are available to external developers with minimal restrictions, fostering innovation by allowing third parties to build new applications on top of existing services. In contrast, **Private APIs** (or Internal APIs) are used exclusively within an organization to connect its own systems and data, improving internal productivity and efficiency. A middle ground is occupied by **Partner APIs**, which are accessible only to specific, authorized business partners to enable strategic collaborations.

The widespread adoption of the RESTful style is a direct result of its architectural advantages. Its stateless, client-server design is the foundation for its **scalability and flexibility**, making it the ideal choice for modern paradigms like **microservices architectures**, where RESTful APIs act as the lightweight glue that binds small, independent services together into a cohesive application. This architectural elegance simplifies communication, reduces development complexity, and ultimately accelerates innovation by allowing developers to efficiently integrate diverse systems and services.

However, this very ubiquity and power also make RESTful APIs a prime target for security threats. Because APIs serve as the direct gateway to an application's most critical functions and sensitive data, a single vulnerability can lead to catastrophic consequences—from massive data breaches to complete service disruption. This report embarks on a comprehensive exploration of RESTful API security, designed to bridge the gap between foundational architectural theory and the cutting-edge of automated testing. To ground our discussion in a real-world context, we begin by documenting the design and implementation of a practical Patient Management API. This hands-on case study will serve as a concrete foundation for the advanced concepts that follow.

With this practical foundation in place, we will then pivot to the forefront of security research, embarking on a detailed literature survey of advanced techniques for automated

black-box testing and intelligent vulnerability detection. We have trace the evolution of testing strategies, from static, dependency-based models to adaptive, reinforcement learning-based approaches that can learn an API's hidden logic. Ultimately, this report aims to synthesize practical implementation with state-of-the-art research to offer a robust framework for developing and securing RESTful APIs, championing a modern, proactive approach where security is not an afterthought, but an integral part of the entire development lifecycle.

Chapter 2

REST API Architecture

A secure and robust RESTful API begins with a strong understanding of its underlying architecture. As established in our practical implementation of the Patient Management API and confirmed by our literature survey, the REpresentational State Transfer (REST) architectural style is not a strict protocol but a design philosophy that leverages standard web protocols, primarily HTTP. This section outlines the core architectural components that define a RESTful API, from the fundamental client-server model to the role of a modern API Gateway.

2.1 The Client-Server Model

The foundational principle of REST is the separation of concerns between the client and the server. The client, or "consumer," is responsible for the user interface and user experience. This can be any application, such as a mobile app on Android or iPhone, a desktop application, or a web browser, as illustrated in Figure 2.1. The server, or "service provider," is responsible for processing requests, applying business logic, and managing data storage.

This decoupling is critical because it allows the client and server to evolve independently. The client does not need to know anything about the server's internal logic or database structure; it only needs to know the endpoint (URI) to send its request to. Similarly, the server can be updated or even completely replaced without affecting the client, as long as the API's interface remains consistent.

REST – Architecture

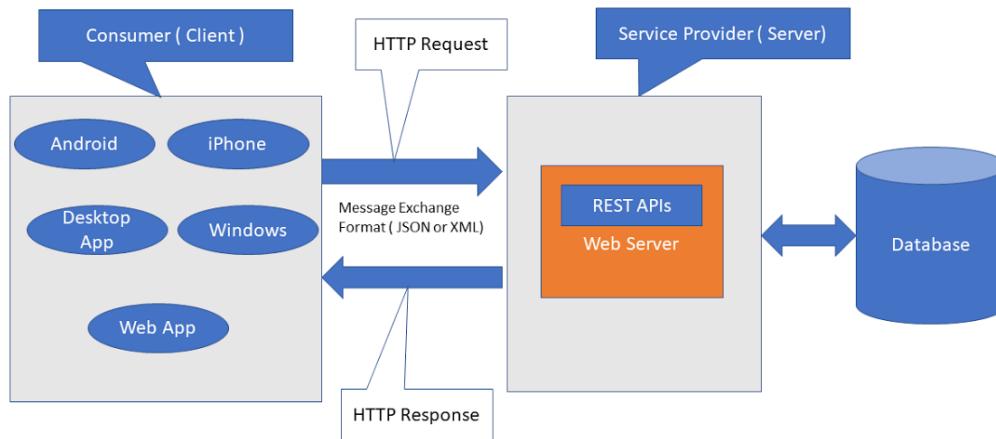


Figure 2.1: The REST Client-Server Architecture, showing various clients interacting with a central server.

2.2 Communication Flow and Core Principles

Communication in a RESTful architecture is governed by two core principles: statelessness and a uniform interface. **Statelessness** means that each request from a client to a server must contain all the information needed to understand and process it. The server does not store any session state from one request to the next, which enhances scalability and reliability.

The **uniform interface** simplifies the architecture by providing a standardized way to interact with resources. This is primarily achieved through standard HTTP methods. As shown in Figure 2.2, a client sends a REST request, which consists of an HTTP method (like GET, POST, PUT, or DELETE) and a URI that identifies the target resource. The server processes this request and sends back a REST response, which is a representation of the resource, typically in a format like JSON or XML.

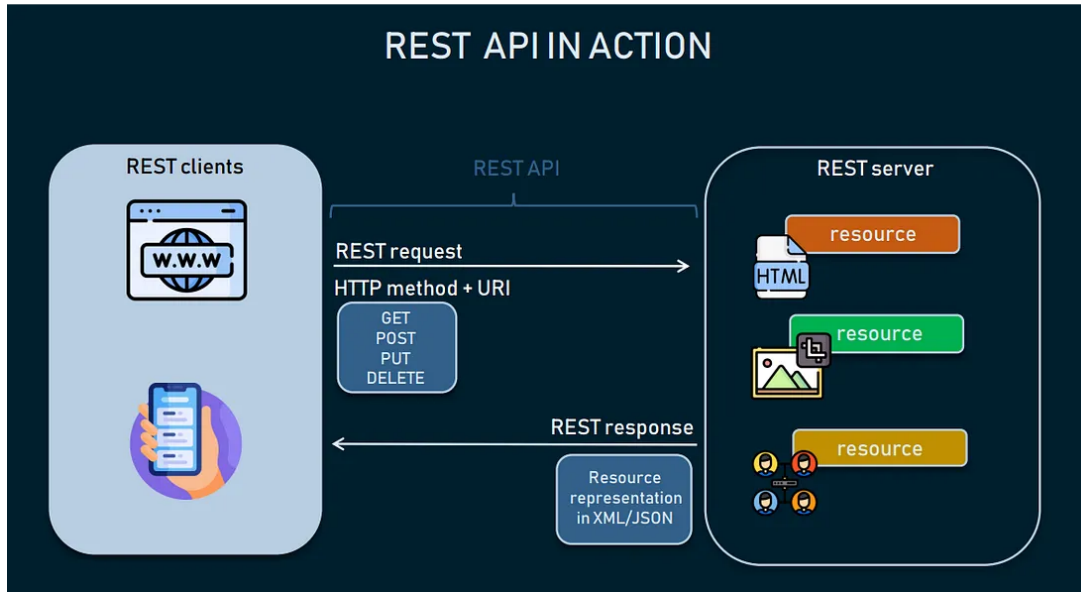


Figure 2.2: The REST API communication flow, illustrating the request-response cycle between client and server.

2.3 The Role of the API Gateway

In modern, complex applications, especially those built on a microservices architecture, a direct client-to-server communication model can become inefficient and insecure. To address this, an **API Gateway** is often introduced as an intermediary layer. An API Gateway acts as a single, unified entry point for all client requests, routing them to the appropriate backend services while handling various cross-cutting concerns.[1, 2]

As depicted in Figure 2.3, an API Gateway is a powerful component that orchestrates the entire request lifecycle. When a client app sends an HTTP request, the gateway performs a series of critical functions:

- **Validation and Verification (Steps 2 & 3):** It first performs parameter validations to ensure the request is well-formed and checks it against whitelists for initial security screening.[2]
- **Security Enforcement (Steps 4 & 5):** The gateway serves as a centralized security checkpoint, handling Authentication and Authorization (AuthN/AuthZ) to verify the client's identity and permissions. It also enforces **rate limiting** to pro-

test backend services from being overwhelmed by too many requests, a key defense against DDoS attacks.[3, 4]

- **Routing and Service Discovery (Steps 6 & 7):** Based on the request, the gateway performs routing to direct the request to the correct backend microservice. It uses service discovery to locate the appropriate service instance dynamically.[5]
- **Request Transformation (Step 8):** Before forwarding the request, the gateway can perform transformations, such as converting protocols or modifying the request payload to match what the backend service expects.[2]
- **Centralized Management:** The gateway also centralizes other essential functions like logging, caching, and error handling, which simplifies the logic of the individual microservices.[6]

By offloading these responsibilities, the API Gateway allows backend services, like the ones we might test with tools from our literature survey, to focus purely on their core business logic. This not only enhances security and performance but also simplifies the overall system architecture.[7]

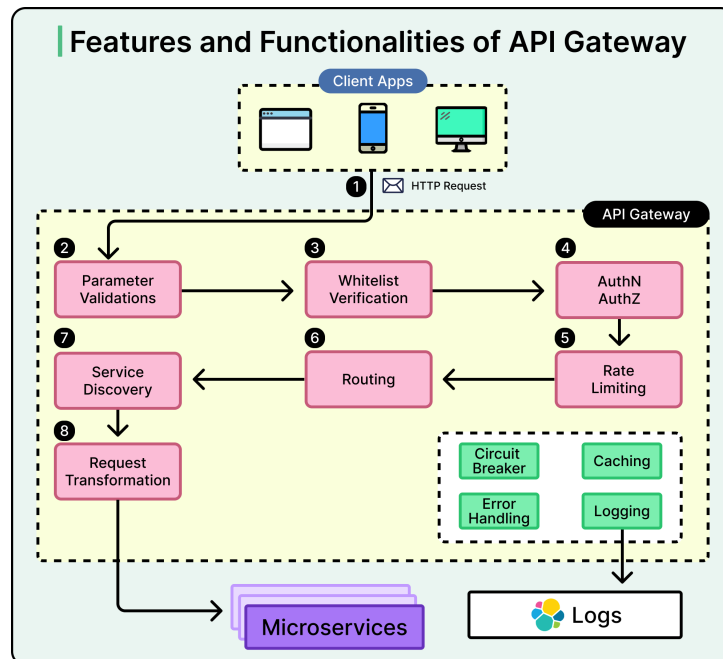


Figure 2.3: Features and functionalities of a modern API Gateway, acting as an intermediary between clients and backend microservices.

Chapter 3

Authentication and Authorization

In API security, authentication and authorization are the two most fundamental gatekeepers. While they work in tandem, they serve distinct purposes. **Authentication** is the process of proving you are who you say you are—like swiping an ID card at a door to verify your identity. **Authorization** happens next; it determines what you are allowed to do once you’re inside—whether your ID card grants you access to your office or the CEO’s suite. Authentication answers, “Who are you?”, while authorization answers, “What are you allowed to do?”.

For a RESTful API, these layers are essential for protecting sensitive data. As our literature survey and practical implementation have underscored, choosing the right strategy has profound implications for an API’s scalability, security, and adherence to core REST principles.

3.1 Session-Based Authentication: The Traditional, Stateful Approach

Session-based authentication is a traditional, stateful approach common in monolithic web applications. When a user logs in, the server creates a “session” record in its own memory or database and sends a unique session ID back to the client, usually in a cookie. For every subsequent request, the browser sends this ID back, and the server uses it to look up the user’s information.

While straightforward, this method’s primary drawback is that it is **stateful**. The server must maintain the state of every active user, which directly violates the stateless principle

of REST. This can introduce scalability challenges in a large, distributed system, as it requires either routing a user to the same server every time or implementing a complex shared session store. For this reason, it is generally considered a less ideal fit for modern, stateless RESTful APIs.

3.2 JWT-Based Authentication: The Modern, Stateless Approach

JSON Web Token (JWT)-based authentication is the de facto standard for securing modern, stateless RESTful APIs. A JWT is a compact, self-contained digital passport. When a user logs in, the server generates a JWT that contains user information (like their ID and roles) and a cryptographic signature to prevent tampering.

This token is sent to the client, which stores it and includes it in the ‘Authorization’ header of every subsequent request. The server can then verify the token’s signature to confirm its authenticity and trust the information within it, without needing to look up a session in a database.

The most significant advantage of this approach is that it is completely **stateless**. All necessary information is encapsulated within the token itself. This perfectly aligns with the core REST principle of statelessness and makes JWTs incredibly scalable, especially for distributed systems and microservices architectures.

3.3 Summary of Authentication Methods

The choice between session-based and JWT-based authentication depends largely on the application’s architecture. The following table provides a side-by-side comparison.

Aspect	Session-Based Authentication	JWT-Based Authentication
State Management	Stateful. The server must store and manage session data for every active user.	Stateless. All necessary user information is contained within the token itself. The server does not need to store session state.
Scalability	Less scalable. Requires sticky sessions or a shared session store in a distributed environment, which adds complexity.	Highly scalable. Any server instance can validate a token, making it ideal for microservices and load-balanced systems.
Token/ID Storage	The server stores the full session object and sends only a small session ID to the client (in a cookie).	The client stores the entire JWT. The server does not need to store anything after issuing the token.
Best Use Case	Traditional, monolithic web applications where the server and client are tightly coupled.	Modern web applications, mobile apps, and distributed systems, especially those using RESTful APIs.
Adherence to REST Principles	Violates the core REST principle of statelessness.	Fully adheres to the statelessness principle of REST, making it the architecturally preferred choice.

Table 3.1: Comparison of Session-Based and JWT-Based Authentication

Chapter 4

Case Study: Patient Management API Implementation

To bridge the gap between the theoretical concepts explored in the literature survey and their real-world application, a fully functional Patient Management RESTful API was designed, developed, and deployed. This project serves as a concrete and tangible case study, demonstrating the core principles of REST architecture in action. It provides a practical system against which the advanced testing and security methodologies from the surveyed literature can be conceptualized and understood. This chapter provides a comprehensive walkthrough of the technical implementation of this API, from the deliberate choices in its technology stack to its key features and live, publicly accessible deployment.

4.1 Technology Stack

The Patient Management API was built using a modern, high-performance Python-based technology stack. Each component was chosen specifically for its ability to facilitate rapid development, ensure robustness, and provide excellent developer tooling, which are critical factors for building and maintaining high-quality APIs.

- **Python:** As the core programming language, Python was selected for its clean syntax, extensive libraries, and strong community support, making it an ideal choice for backend development.
- **FastAPI:** This modern, high-performance web framework was the cornerstone of

the project. It was selected for its exceptional speed and native support for asynchronous operations. A crucial advantage of FastAPI is its ability to automatically generate interactive API documentation and a machine-readable OpenAPI Specification directly from the Python code, which dramatically accelerates both development and testing cycles.

- **Pydantic:** In a dynamically typed language like Python, ensuring data integrity is paramount. Pydantic was instrumental in this regard, providing a powerful data validation layer that uses standard Python type hints. It enforces a strict data schema at runtime, guaranteeing that all incoming requests conform to the expected data models and preventing malformed or invalid data from ever reaching the application's business logic.
- **Uvicorn:** To serve the application, Uvicorn, a lightning-fast ASGI (Asynchronous Server Gateway Interface) server, was used. Its high performance makes it a perfect match for running asynchronous frameworks like FastAPI.
- **Render:** For deployment, the Render cloud platform was chosen. Its ease of use and seamless integration with code repositories allowed for the straightforward deployment of the API, making it publicly accessible via a stable URL for live demonstration and testing.

4.2 API Endpoint Specification

The API provides a comprehensive and intuitive set of endpoints for managing patient records. The functionality of each endpoint, along with its expected success and failure responses, is formally documented in Table 4.1. This clear and explicit contract is for both client developers to integrate with the API and for the automated testing tools, as it allows them to understand the outcomes of their requests and validate the API.

Endpoint & Method	Functional Description	Success Code	Common Failure Codes
POST /create	Creates a new patient record in the system. The request body must conform to the Patient schema.	201 Created	400 Bad Request (if patient ID already exists or input data is invalid).
GET /view	Fetches a list of all patient records currently stored in the database.	200 OK	500 Server Error (if the data store is inaccessible or corrupt).
GET /patient/{id}	Fetches the details of a single patient by their unique ID provided in the path.	200 OK	404 Not Found (if no patient with the given ID exists).
PUT /edit/{id}	Updates the information of an existing patient. Only the fields provided in the request body are modified.	200 OK	404 Not Found (if no patient with the given ID exists). 400 Bad Request (if update data is invalid).
DELETE /delete/{id}	Deletes a patient record from the system based on the unique ID provided in the path.	200 OK	404 Not Found (if no patient with the given ID exists).
GET /sort	Returns a list of patients sorted by a specified attribute (height, weight, or bmi) and order (asc, desc).	200 OK	400 Bad Request (if the sort_by or order query parameters are invalid).

Table 4.1: Functional Specification of the Patient API Endpoints

4.3 Key Features and Functionality

Beyond the standard CRUD operations detailed above, the API implements several key features designed to ensure data integrity, provide useful functionality, and enhance the developer experience.

- **Robust Data Validation:** Leveraging Pydantic, the API enforces strict, server-side validation rules on all incoming data. This serves as a critical first line of defense against data corruption. For example, a patient's age must be a valid integer between 1 and 119, their gender must be one of the predefined literals ('male', 'female', or 'others'), and their height and weight must be positive numbers. Any request that violates these rules is automatically rejected with a descriptive '400 Bad Request' error, ensuring that only clean, valid data is processed.
- **Integrated Business Logic:** The API goes beyond simple data storage by integrating business logic directly into its data model. The Body Mass Index (BMI) is automatically calculated from the patient's height and weight, and a corresponding health 'verdict' ('Underweight', 'Normal', or 'Obese') is assigned based on this BMI value. This logic is executed automatically whenever a patient record is created or updated, ensuring that these derived health metrics are always consistent and up-to-date without requiring any extra work from the client.
- **Automatic Interactive Documentation:** A standout feature provided by the FastAPI framework is the automatic generation of interactive API documentation via Swagger UI. This creates a user-friendly web interface, as shown in Figure 4.1, that serves as a live, interactive playground for the API. It allows developers and testers to visualize and interact with all API endpoints directly from their browser, view schemas, and execute test requests without writing a single line of code. This significantly speeds up development, testing, and integration cycles.

4.4 Deployment and Live Demonstration

The Patient Management API is not just a theoretical or local project; it has been successfully deployed to the Render cloud platform and is publicly accessible. This live deployment transforms it from a development exercise into a real-world, testable service that can be interacted with by any HTTP client from anywhere in the world, demonstrating its readiness and stability.

Live Demo Link: <https://patient-api-zo4v.onrender.com/docs>

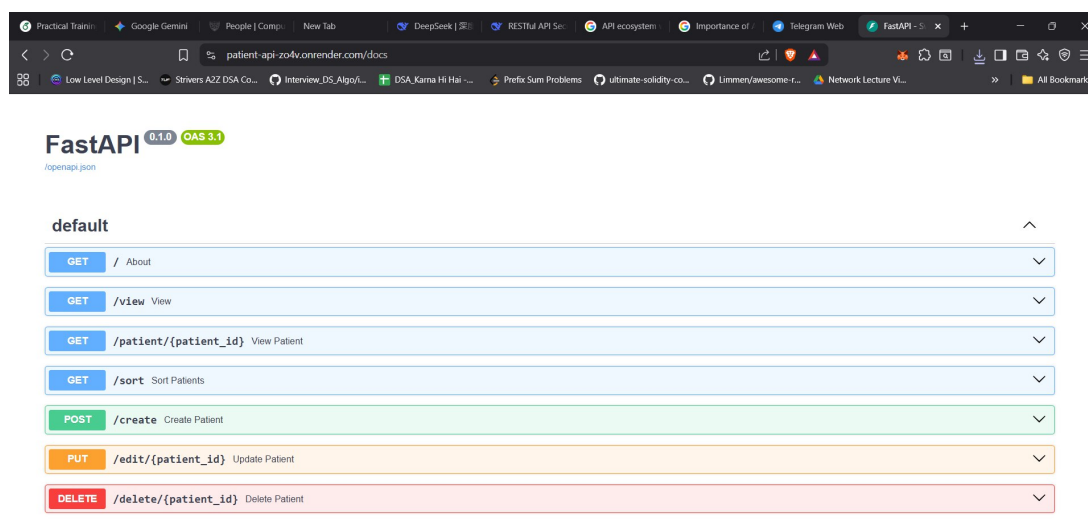


Figure 4.1: The interactive Swagger UI documentation, automatically generated by FastAPI, displaying all available endpoints for the Patient API.

Chapter 5

Common API Security Threats

The open and accessible nature of RESTful APIs makes them a prime target for malicious attacks. A robust security strategy must anticipate and mitigate these threats from the design phase, as an API often serves as the digital front door to an application's most valuable data and logic.

5.1 Denial of Service Attacks

A Distributed Denial of Service (DDoS) attack is a malicious attempt to make a service unavailable by overwhelming it with a flood of internet traffic, typically from a network of compromised computers known as a botnet. Key mitigation strategies include:

- **Rate Limiting:** This is a crucial first line of defense that controls the number of requests a user or IP address can make to an API within a specific timeframe. If the rate exceeds a predefined threshold, the server throttles or blocks further requests from that source.
- **Traffic Filtering:** More advanced strategies use specialized cloud services or network hardware to analyze and filter out malicious traffic patterns before they ever reach the application server.

5.2 The OWASP API Security Top 10

The OWASP Foundation provides an essential list of the top 10 most critical security risks for APIs, which serves as a foundational guide for developers. Understanding and

addressing these risks is fundamental to building secure APIs. We will briefly explore three of the most critical threats.

- **Broken Object Level Authorization (BOLA):** Ranked as the number one threat, BOLA occurs when an API fails to verify that a user has the right to access the specific object they are requesting. For example, in our Patient Management API, a user authenticated as patient P002 could simply change the ID in a request to GET /patient/P001 and, if BOLA is present, illegitimately access another patient's sensitive data.

Mitigation: Implement strict, explicit authorization checks for every request that accesses a resource. The server must verify that the authenticated user's ID matches the owner ID of the requested resource.

- **Broken User Authentication:** It flaws in the authentication mechanism itself, such as allowing weak passwords or improperly handling authentication tokens. If an attacker can bypass or compromise the authentication process, they can impersonate users and gain unauthorized access to the system.

Mitigation: A robust authentication system requires enforcing strong password policies, implementing multi-factor authentication (MFA), and using secure, industry-standard methods for token generation and validation.

- **Excessive Data Exposure:** It occurs when an API returns more data in its response than the client application actually needs. Even if the extra data is not displayed on the user interface, an attacker can intercept the raw API response and access any sensitive information. **Mitigation:** API endpoints should be carefully designed to return only the minimum amount of data required for a specific action, following the principle of least privilege.

Chapter 6

LITERATURE SURVEY

6.1 Survey Table

Table 6.1: Literature Survey of Key Research Papers

Sl No.	Paper	Objective	Approach	Limitation
1	Atlidakis et al. (2019) <i>RESTler: Stateful REST API Fuzzing</i>	Automatically test cloud services via REST APIs by generating and executing stateful sequences of requests.	Infers producer-consumer dependencies from OpenAPI spec and uses dynamic feedback from responses to generate new tests and prune the search space.	Primarily detects bugs with 500 HTTP status codes, doesn't detect other vulnerabilities (e.g., information exposure).
2	Liu et al. (2022) <i>MOREST: Model-based RESTful API Testing...</i>	Generate meaningful API call sequences for black-box testing, overcoming inconsistencies in API specifications.	Builds and maintains a dynamically updating RESTful-service Property Graph (RPG) to model behaviors and guide call sequence generation.	Model quality depends on the OpenAPI specification's quality, which can still be ambiguous or poorly written.

3	Corradini et al. (2024) <i>DeepREST: Automated Test Case Generation...</i>	Automatically generate effective black-box test cases for REST APIs by learning implicit constraints.	Leverages deep reinforcement learning to uncover hidden API constraints, guiding an agent in exploring API states and generating accurate input data.	High computational complexity and intensive training required for the DRL model.
4	Deng et al. (2023) <i>NAUTILUS: Automated RESTful API Vulnerability Detection</i>	Uncover RESTful API vulnerabilities, particularly multi-API vulnerabilities, that require specific sequences of operations to trigger.	Uses novel specification annotations to guide the generation of meaningful operation sequences and injects vulnerability-specific payloads.	Vulnerability detection depends on a predefined payload dictionary and may miss vulnerabilities requiring case-specific payloads.
5	Kim et al. (2023) <i>Adaptive REST API Testing with Reinforcement Learning</i>	Incorporate reinforcement learning to prioritize operations and parameters during exploration for adaptive testing.	A Q-learning-based approach dynamically analyzes request/response data and uses a sampling strategy to efficiently process API feedback.	May struggle with semantic parameters that require inputs in specific, non-random formats (e.g., email addresses).
6	Davide Corradini et al. (2010) <i>Automated Black-Box Testing of Nominal and Error Scenarios in RESTful APIs</i>	Automatically generate black-box test cases for REST APIs by modeling data dependencies.	Uses an Operation Dependency Graph (ODG) to dynamically decide the order to test operations, reusing data from outputs as inputs for subsequent operations.	The ODG's contribution can be negligible or even detrimental when no data dependencies are available in the graph.

Summary

This summary provides a concise overview of the key literature reviewed, highlighting the contributions and limitations of each work in the context of automated RESTful API security testing.

- **RESTler: Stateful REST API Fuzzing (Atlidakis et al., 2019):** This paper introduced **RESTler**, a stateful fuzzer that performs a lightweight static analysis of an OpenAPI specification to infer **producer-consumer dependencies**. This approach allows it to generate sequences of requests that explore deeper service logic, uncovering bugs that simple, non-sequential testing would miss. Its key strength is the combination of static analysis with dynamic feedback to prune the vast search space and focus on effective test sequences.
- **MOREST: Model-based RESTful API Testing... (Liu et al., 2022):** The MOREST framework proposes a sophisticated **model-based black-box testing** technique using a dynamically updating **RESTful-service Property Graph (RPG)**. This graph models dependencies between operations and schemas, enabling the generation of longer, more meaningful call sequences that adapt to execution feedback. This approach helps to overcome ambiguities in API specifications and can find bugs that require complex, multi-step scenarios.
- **DeepREST: Automated Test Case Generation... (Corradini et al., 2024):** DeepREST, a novel black-box approach, leverages **deep reinforcement learning (DRL)** to uncover implicit API constraints that are not explicitly defined in the OpenAPI specification. A DRL agent learns an effective order of operations and the best strategies for generating input values, going beyond what can be statically inferred. This allows it to find bugs related to the API’s business logic and complex state transitions.

- **NAUTILUS: Automated RESTful API Vulnerability Detection (Deng et al., 2023):** NAUTILUS is an automated vulnerability detection tool that uses a novel **specification annotation strategy** to uncover RESTful API vulnerabilities. It focuses specifically on **multi-API vulnerabilities** that require specific sequences of operations to trigger. NAUTILUS extends the OpenAPI specification with annotations that encode proper operation relations and parameter generation strategies, allowing it to generate meaningful attack sequences and use specific payload dictionaries to find security flaws like SQL injection, XSS, and privilege escalation.
- **Davide Corradini, et al. Automated Black-Box Testing of Nominal and Error Scenarios in RESTful APIs (Kim et al., 2023):** This paper presents **ARAT-RL**, an adaptive testing technique that uses **Q-learning** to prioritize the exploration of operations and parameters. It dynamically constructs key-value pairs from both request and response data, which helps to deal with incomplete or imprecise response schemas. The approach’s primary goal is to increase code coverage and fault detection by intelligently navigating the test space based on learning from past interactions.
- **Davide Corradini, et al. Automated Black-Box Testing of Nominal and Error Scenarios in RESTful APIs:** This work presents **RESTTESTGEN**, a black-box approach that automatically generates test cases for both nominal and error scenarios based on an API’s OpenAPI specification. Its core contribution is the **Operation Dependency Graph (ODG)**, which models data dependencies between operations to determine an effective testing order. By testing producer operations before consumer operations, it can reuse data from outputs as inputs for subsequent calls, addressing the challenge of generating valid input values.

Chapter 7

RESEARCH GAPS

7.1 Identified Research Gaps

Despite significant advancements, our literature survey reveals several critical gaps in the automated black-box testing of REST APIs. A primary challenge is **learning implicit business logic**, as most tools rely on explicit dependencies in the OpenAPI specification and struggle with undocumented state transitions. There is also a need for more **comprehensive vulnerability detection** frameworks that can execute the multi-step, stateful attacks required to find security flaws, rather than just functional bugs.

Furthermore, the effectiveness of most tools is hindered by their inability to handle **incomplete or poorly-written specifications**, a common issue in real-world projects. Advanced techniques like deep reinforcement learning introduce **scalability and performance trade-offs**, as they can be computationally intensive.

7.2 Strategies for Addressing the Gaps

To bridge these gaps, future research can focus on several promising strategies. The development of **hybrid testing frameworks** that combine the strengths of different approaches—such as a fuzzer’s ability to mutate requests with a reinforcement learning agent’s strategic sequence generation—could offer a more powerful and efficient solution.

Additionally, tools could be enhanced with **advanced machine learning models**, such as Large Language Models (LLMs), to generate more context-aware inputs and better

understand API purposes. To overcome flawed documentation, tools could incorporate **self-correcting logic**, allowing them to learn and adapt when an API's actual behavior deviates from its specification. Finally, a focus on **optimized algorithms** and parallel processing could help mitigate the performance overhead of more advanced testing techniques, making them more practical for large-scale APIs.

Chapter 8

Conclusion

This report has provided a comprehensive journey through the landscape of RESTful API security, beginning with the foundational architectural principles and culminating in a survey of advanced, research-driven testing methodologies. By grounding this exploration in the practical implementation of a Patient Management API, we have demonstrated not only the core components of a modern API but also the tangible challenges that arise in securing and testing such a system.

The literature survey confirmed that the field of automated testing is rapidly evolving. State-of-the-art tools have moved beyond simple, stateless checks to embrace complex, stateful fuzzing and even reinforcement learning to uncover deep-seated bugs. However, this survey also illuminated several critical research gaps that prevent the full realization of truly autonomous and intelligent API testing. The reliance on often-flawed specifications, the difficulty in learning implicit business logic, and the challenge of generating semantically valid data remain significant hurdles for the current generation of tools.

Chapter 9

Future Work

The path forward for automated REST API security testing lies in directly addressing the gaps identified in this report. The development of **hybrid testing frameworks** that combine the strengths of different approaches is a key area for exploration. A system that integrates a fuzzer’s ability to mutate requests with a reinforcement learning agent’s strategic sequence generation could offer a more powerful and efficient testing pipeline that is greater than the sum of its parts.

Additionally, testing tools could be significantly enhanced with **advanced machine learning models**, such as Large Language Models (LLMs), to generate more context-aware inputs and better understand an API’s purpose. To overcome the persistent problem of flawed documentation, future tools could incorporate **self-correcting logic**, allowing them to learn and adapt when an API’s actual behavior deviates from its specification. Finally, a focus on **optimized algorithms** and parallel processing will be crucial to mitigate the performance overhead of these advanced techniques, making them more practical and scalable for testing the large, complex APIs that power modern applications.

REFERENCES

- [1] Atlidakis, V., Godefroid, P., & Polishchuk, M. (2019). RESTler: Stateful REST API Fuzzing. In *2019 IEEE/ACM 41st International Conference on Software Engineering (ICSE)*. IEEE.
- [2] Liu, Y., Li, Y., Deng, G., Liu, Y., Wan, R., Wu, R., Ji, D., Xu, S., & Bao, M. (2022). MOREST: Model-based RESTful API Testing with Execution Feedback. In *2022 IEEE/ACM 44th International Conference on Software Engineering (ICSE)*. IEEE.
- [3] Corradini, D., Montolli, Z., Pasqua, M., & Ceccato, M. (2024). DeepREST: Automated Test Case Generation for REST APIs Exploiting Deep Reinforcement Learning. In *39th IEEE/ACM International Conference on Automated Software Engineering (ASE)*. ACM.
- [4] Deng, G., Zhang, Z., Li, Y., Liu, Y., Zhang, T., Liu, Y., Yu, G., & Wang, D. (2023). NAUTILUS: Automated RESTful API Vulnerability Detection. In *2023 IEEE/ACM International Conference on Automated Software Engineering (ASE)*.
- [5] Kim, M., Sinha, S., & Orso, A. (2023). Adaptive REST API Testing with Reinforcement Learning. *arXiv preprint arXiv:2309.04583*.
- [6] Viglianisi, E., Dallago, M., & Ceccato, M. (2020). RESTTESTGEN: Automated black-box testing of RESTful APIs. In *2020 IEEE 13th International Conference on Software Testing, Validation and Verification (ICST)*. IEEE.

- [7] Fielding, R. T. (2000). *Architectural Styles and the Design of Network-based Software Architectures*. Doctoral Dissertation, University of California, Irvine.
- [8] The OWASP Foundation. (2023). *API Security Top 10*. OWASP Project. Retrieved from <https://owasp.org/www-project-api-security/>
- [9] Richardson, C. (2018). *Microservices Patterns: With examples in Java*. Manning Publications. (Provides foundational concepts on the API Gateway pattern).
- [10] FastAPI. (2024). *Security - First Steps*. FastAPI Official Documentation. Retrieved from <https://fastapi.tiangolo.com/tutorial/security/>
- [11] Postman. (2024). *API Security Best Practices*. Postman Learning Center. Retrieved from <https://learning.postman.com/docs/sending-requests/authorization/>